AI Detection Report for My normal flow is to have these long, meandering s...

Aug 29, 2026 | 1,307 Words | Pangram 4.0

Summary

We believe that this entire text is human-written.

Authorship Breakdown

0 %	of the document	AI Generated
0 %	of the document	Moderately AI Assisted
0 %	of the document	Lightly AI Assisted
100 %	of the document	Human Written

Authorship Progression

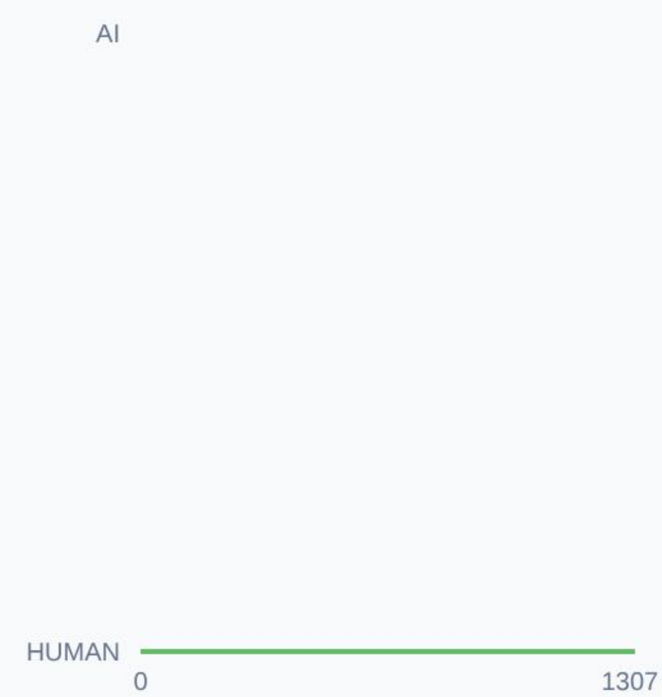

Analyzed Text

Words that are underlined indicate AI phrases, while the numbers represent how frequently they are likely to appear in AI writing. These elements are not used to determine the results but are displayed separately as patterns commonly found in AI-generated text.

Human Written | 1307 Words | High Confidence

My normal flow is to have these long, meandering streams of consciousness — this is one of them — and then take the transcript and ask an AI to clean it up. Which is necessary, because the AI makes typos when it transcribes. So in some sense, if you are simply using AI to make the transcript's fidelity to your words higher, that seems perfectly kosher. In some sense you are decreasing the AI, because rather than allowing it to make mistakes in your words, you are allowing it to fix them. And I think AI is pretty good at this. You tell it: go through my transcript, look for transcription typos, and use context and the broad body of what I'm saying to fill in the words you think I intended. To me, that seems totally fine. Now, the second step is that you can then have Claude try to package your arguments into something more coherent. And you can still say, hey, use my own words. But at this point, when the structure of the argument is sort of outsourced, you do lose the human soul of your writing. So I live in two worlds. In one, I'll record a voice note and allow Claude to massage it into a more structured argument. In the other, I'm just rapidly stream-of-consciousness typing on Slack or iMessage. And what I'm lacking, what I haven't done since college, to be honest, is a coherent long form. Not even long form — not an essay, but a piece of personal writing that is wholly human, that has no element of the machine. So that's what I'm going to attempt to do here. I'm going to take this voice note, have Claude fix the transcription errors, and then ultimately formulate it into a piece of writing myself. What's most interesting now is that we've lived in this world where AI

Human Written | 1307 Words | High Confidence

detection was, frankly, a farce — false positives were way too high — and now we live in a world where you actually can do it. With Pangram, it is remarkably difficult to fool. And this is actually quite an interesting update: no matter how smart the model is, GPT-5, Claude 5.6, there seems to be something. You can feel it. This is kind of a meta point about why I'm doing this. You can feel what a Claude-written essay is shaped like. There is some intuitive pattern recognition where within a sentence or two I can be like, oh, this is not a human essay. And this goes down to even just if you ask an AI to list numbers, or random words. That list of words or numbers can be detected as AI by the systems. So there is fundamentally a patterning going on. And if you think about it, for coding the patterning is good, because there is such a thing as good code. Style doesn't really matter — I mean, style does matter, I suppose. You can have code that works, but it's overly defensive and bloated, which is kind of a critique of GPT-5.6. But in writing, there isn't some single universal best way to write. There are tropes, but then the tropes sort of collapse into a caricature of themselves. If you use all the best writing tropes, which is maybe what Claude does, you have something that is better, maybe, than 80% of writers in terms of technical merit, but ultimately has some soullessness in it. And this goes back to the Apollonian and the Dionysian. It's strange, right? You would think that a probabilistic machine would have a sort of raw, vibrant, fertile chaos. And it does. The base model clearly has this — you can see almost the frenzy of their style. You can even see it in some of these Hugging Face hacks, where you look at the internal reasoning chain and there is something sort of wild, or inventive. And somehow, in these post-trained models, when it comes to writing — how ironic is it that of all the things a language model can do, the thing it's actually worst at is writing? It's strange. But again, it's not strange if you realize that it's all about verifiability, and the idea of, is there a single best way to do this? That's where taste comes in. Taste, in some sense, is that there are not a hundred best ways to solve a coding problem. There are objective measures as to which is the best code — the thing that does what you want, most reliably, with as few lines of code as possible. That's maybe a simple way to put it. But for writing, what you want in some deep sense is almost a parasocial relationship with the author's mind. And the real betrayal of reading something that was written by AI is that you have this uncanny sense that you thought you were going to commune with the mind of the writer, the voice of the writer, the consciousness of the writer, in the most visceral way you can access consciousness, which is words, and the choice of words. And when you feel the pattern of Claude slop, you feel betrayed. Because the thing you really wanted was not the content — I don't think it's chiefly the content. You wanted the form. You wanted the soul of the writer. And instead you got, essentially, the selection. It's still valuable in a sense. The writer is still curating content for you. But you no longer have access to the raw — I mean, look, you never access anyone's raw consciousness. And that goes to this deeper point, which is that we all live in these hermetically sealed bubbles. We can never see outside of our own first-personhood. There's the philosophical zombie problem: do we even know that there's any light behind the eyes inside the minds of anyone else, or are they all just zombies? And obviously conversation and writing, the trading of words, is one way in which you can transcend the boundaries of self. So I think that's almost the spiritual value of writing and reading. And when you pass everything through the layer of Claude, you divorce form and content, and what you realize is that no one really cared about content. They cared about the form. And look, this is maybe different for fiction and nonfiction. Nonfiction is more information retrieval in some sense. If you're reading a self-help book, or a technical book on some matter, it really is: is the information correct, and is it written in a legible manner? But the best nonfiction is more

Human Written | 1307 Words | High Confidence

than that. The best nonfiction is story, it's character. It's weaving the true facts on the ground into the structure of a narrative in a way that makes it maximally compelling while also teaching you. And that requires the voice of the writer. If you want to just read an encyclopedia page, then AI is perfectly good. So, yeah, I think that wraps up the thought. The exercise now is to have Claude fix this, and then do it on my own. But isn't this something interesting? Can this leverage the raw thought itself? I can imagine an interesting thing, which is almost like a blog. It starts with the voice note. You include the audio, you include the transcript, then you include the Claude-cleaned-up transcript. At the very top you would include your article, and maybe you also include a Claude article, so you can sort of compare the difference. That might be an interesting format. I like that idea.